

¹ OMEinsumContractionOrders: A Julia package for tensor network contraction order optimization

³ Jin-Guo Liu  ^{1*}, Xuanzhao Gao^{2*}, and Richard Samuelson^{3*}

⁴ 1 Hong Kong University of Science and Technology (Guangzhou) 2 Center of Computational
⁵ Mathematics, Flatiron Institute 3 * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

⁶ Statement of need

⁷ OMEinsumContractionOrders (One More Einsum Contraction Orders, or OMECO) is a Julia ⁸ package ([Bezanson et al., 2012](#)) that implements state-of-the-art algorithms for optimizing ⁹ tensor network contraction orders. This paper presents its key features, integration with the ¹⁰ Julia ecosystem, and performance benchmarks.

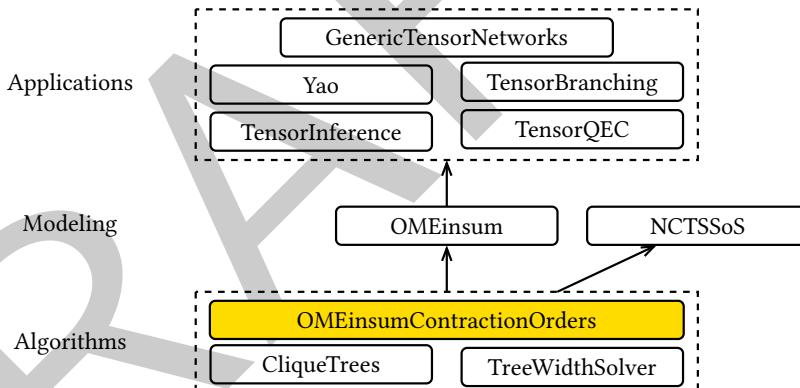


Figure 1: The ecosystem built around OMEinsumContractionOrders and its dependencies. OMECO serves as a core component of the tensor network contractor OMEinsum, which powers applications including Yao (quantum simulation), TensorQEC (quantum error correction), TensorInference (probabilistic inference), GenericTensorNetworks and TensorBranching (combinatorial optimization).

¹¹ A *tensor network* is a mathematical framework that represents multilinear algebra operations as ¹² graphical structures, where tensors are nodes and shared indices are edges. This diagrammatic ¹³ approach transforms complex high-dimensional contractions into visual networks that expose ¹⁴ underlying computational structure.

¹⁵ The framework has remarkable universality across diverse domains: *einsum* notation ([Harris et al., 2020](#)) in numerical computing, *factor graphs* ([Bishop & Nasrabadi, 2006](#)) in probabilistic ¹⁷ inference, *sum-product networks* in machine learning, and *junction trees* ([Villegas et al., 2023](#)) in graphical models. Tensor networks have enabled breakthroughs in quantum circuit ¹⁹ simulation ([Markov & Shi, 2008](#)), quantum error correction ([Piveteau et al., 2024](#)), neural ²¹ network compression ([Qing et al., 2024](#)), strongly correlated quantum materials ([Haegeman et al., 2016](#)), and combinatorial optimization problems ([Liu et al., 2023](#)). These applications are ²² reflected in the ecosystem built around OMECO, as illustrated in [Figure 1](#).

²³ The computational cost of tensor network contraction depends critically on the *contraction* ²⁴ order—the sequence in which pairwise tensor multiplications are performed. This order

25 can be represented as a binary tree where leaves correspond to input tensors and internal
 26 nodes represent intermediate results. The optimization objective balances multiple complexity
 27 measures through the cost function:

$$\mathcal{L} = w_t \cdot tc + w_s \cdot \max(0, sc - sc_{target}) + w_{rw} \cdot rwc,$$

28 where w_t , w_s , and w_{rw} represent weights for time complexity (tc), space complexity (sc), and
 29 read-write complexity (rwc), respectively. In practice, memory access costs typically dominate
 30 computational costs, motivating $w_{rw} > w_t$. The space complexity penalty activates only when
 31 $sc > sc_{target}$, allowing unconstrained optimization when memory fits within available device
 32 capacity.

33 Finding the optimal contraction order—even when minimizing only time complexity—is NP-
 34 complete (Markov & Shi, 2008). This optimization problem has a deep mathematical connection
 35 to *tree decomposition* (Markov & Shi, 2008) of the tensor network's line graph, where finding
 36 the optimal order corresponds to finding a weighted minimal-width tree decomposition. The
 37 logarithmic time complexity of the bottleneck contraction step equals the largest bag size in
 38 the tree decomposition, while the logarithmic space complexity equals the largest separator
 39 size (vertices shared between adjacent bags).

40 Despite this computational hardness, near-optimal solutions suffice for most practical ap-
 41 plications and can be obtained efficiently through heuristic methods. Modern optimization
 42 algorithms have achieved remarkable scalability, handling tensor networks with over 10^4 tensors
 43 (Gray & Kourtis, 2021; Roa-Villegas et al., 2024).

44 OMECO implements several optimization algorithms with complementary performance charac-
 45 teristics:

Optimizer	Description
GreedyMethod	Fast greedy heuristic with modest solution quality
TreeSA	Reliable simulated annealing optimizer (Kalachev et al., 2021) with high-quality solutions
PathSA	Simulated annealing optimizer for path decomposition
HyperND	Nested dissection algorithm for hypergraphs, requires KaHyPar or Metis
KaHyParBipartite	Graph bipartition method for large tensor networks (Gray & Kourtis, 2021), requires KaHyPar
SABipartite	Simulated annealing bipartition method, pure Julia implementation
ExactTreewidth	Exact algorithm with exponential runtime (Bouchitté & Todinca, 2001), based on TreeWidthSolver
Treewidth	Clique tree elimination methods from CliqueTrees package

46 The algorithms HyperND, Treewidth, and ExactTreewidth operate on the tensor network's line
 47 graph and utilize the CliqueTrees and TreeWidthSolver packages, as illustrated in Figure 1.
 48 Additionally, the PathSA optimizer implements path decomposition by constraining contraction
 49 orders to path graphs, serving as a variant of TreeSA.

50 These methods balance optimization time against solution quality. Figure 2 displays benchmark
 51 results for the Sycamore quantum supremacy circuit, highlighting the Pareto front where
 52 contraction order quality is balanced with optimization runtime.

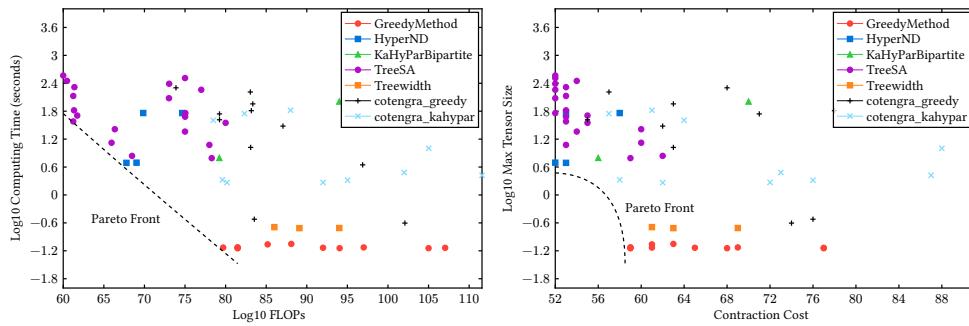


Figure 2: Time complexity (left) and space complexity (right) benchmark results for contraction order optimization on the Sycamore quantum circuit tensor network (Intel Xeon Gold 6226R CPU @ 2.90GHz, single-threaded). The x -axis shows contraction cost, y -axis shows optimization time. Each point represents a different optimizer configuration tested with varying parameters. TreeSA and HyperND achieve the lowest contraction costs, while GreedyMethod offers the fastest optimization time.

53 [JG: TODO: Please also cite CliqueTree paper.]
 54 Optimizers prefixed with `cotengra_` are from the Python package `cotengra` (Gray & Kourtis,
 55 2021); all others are OMECO implementations. For both optimization objectives (minimizing
 56 time and space complexity), OMECO optimizers dominate the Pareto front. Given suffi-
 57 cient optimization time, TreeSA consistently achieves the lowest time and space complexity.
 58 GreedyMethod provides the fastest optimization but yields suboptimal contraction orders, while
 59 HyperND offers a favorable balance between optimization time and solution quality.
 60 OMECO has been integrated into the OMEinsum package and powers several downstream
 61 applications: Yao (Luo et al., 2020) for quantum circuit simulation, GenericTensorNetworks
 62 (Liu et al., 2023) and TensorBranching (TODO: add citation) for combinatorial optimization,
 63 TensorInference (Roa-Villegas & Liu, 2023) for probabilistic inference, and TensorQEC for
 64 quantum error correction. This infrastructure is expected to benefit other applications requiring
 65 tree or path decomposition, such as polynomial optimization (Magron & Wang, 2021).

66 Usage Example

67 OMECO provides two main functions: `optimize_code` for finding optimal contraction orders,
 68 and `slice_code` for trading time complexity for reduced space complexity through the slicing
 69 technique.

70 To demonstrate basic usage, we generate a random 3-regular graph with 100 vertices using the
 71 `Graphs` package, associating each vertex with a binary variable and each edge with a 2×2
 72 tensor.

```
julia> using Graphs: random_regular_graph, edges, vertices

julia> using OMEinsumContractionOrders: EinCode, uniquelabels, contraction_complexity, o

julia> function demo_network(n::Int)
       g = random_regular_graph(n, 3)
       code = EinCode([[e.src, e.dst] for e in edges(g)], Int[])
       sizes = Dict(i=>2 for i in uniquelabels(code))
       tensors = [randn([sizes[index] for index in ix]...) for ix in code.iks]
       return code, tensors, sizes
end
demo_network (generic function with 1 method)
```

```
julia> code, tensors, sizes = demo_network(100);

73 The tensor network topology is represented by an EinCode object with two fields: ixs (a
74 vector of index vectors for each input tensor) and iy (output indices). This structure defines
75 a hypergraph with potentially open edges. Combining this hypergraph with tensor sizes
76 determines the contraction complexity.

julia> contraction_complexity(code, sizes)
Time complexity: 2^100.0
Space complexity: 2^0.0
Read-write complexity: 2^9.231221180711184

77 The return type contains three fields (tc, sc, rwc) for time, space, and read-write complexity.
78 Without optimization, the time complexity is  $2^{100}$ , equivalent to brute-force enumeration.

79 We now use the TreeSA optimizer to find an improved contraction order.

julia> optcode = optimize_code(code, sizes, TreeSA(; score=ScoreFunction(tc_weight=1.0))

julia> cc = contraction_complexity(optcode, sizes)
Time complexity: 2^17.241796993093228
Space complexity: 2^13.0
Read-write complexity: 2^16.360864226366807

80 The optimize_code function takes three arguments: the EinCode object, tensor size dictionary,
81 and optimizer configuration. It returns a NestedEinsum object specifying the contraction tree
82 with three fields: args (child nodes), tensorindex (input tensor index for leaf nodes), and
83 eins (einsum notation for the node). The time complexity  $\approx 2^{17.2}$  is dramatically improved
84 from the original  $2^{100}$ . This result aligns with theory, as the treewidth of a 3-regular graph is
85 approximately upper bounded by  $1/6$  of the number of vertices (Fomin & Høie, 2006). The
86 score keyword argument configures the cost function weights; here we set the read-write
87 weight to  $10 \times$  the time weight, reflecting the higher cost of memory access.

88 Space complexity can be further reduced using slice_code, which implements the slicing
89 technique to trade time for space.

julia> sliced_code = slice_code(optcode, sizes, TreeSASlicer(score=ScoreFunction(sc_targ

julia> sliced_code.slicing
3-element Vector{Int64}:
 14
 76
 60

julia> contraction_complexity(sliced_code, sizes)
Time complexity: 2^17.800899899920303
Space complexity: 2^10.0
Read-write complexity: 2^17.199595668955244

90 The slice_code function takes the NestedEinsum object, tensor sizes, and slicing strat-
91 egy, returning a SlicedEinsum object with two fields: slicing (sliced indices) and eins (a
92 NestedEinsum object). Using TreeSASlicer, we reduce space complexity by  $2^3$  (from  $2^{13}$ 
93 to  $2^{10}$ ) with only a modest increase in time complexity. The resulting SlicedEinsum object
94 maintains the same interface as NestedEinsum for contraction evaluation.

julia> @assert sliced_code(tensors...) ≈ optcode(tensors...)

95 [JG: TODO: Mention the API to convert between contraction graph and treewidth. (Xuan-Zhao
96 fill in), Remove?]
```

⁹⁷ [JG: TODO: Show a plot about using slicing to reduce the space complexity (based on the
⁹⁸ above example). (Xuan-Zhao fill in)]

⁹⁹ Real-world examples demonstrating applications to quantum circuit simulation, combinatorial
¹⁰⁰ optimization, and probabilistic inference are available in the [OMEinsumContractionOrders-](#)
¹⁰¹ [Benchmark](#) repository. Optimizer performance is highly problem-dependent, with no single
¹⁰² algorithm dominating across all metrics and graph topologies.

¹⁰³ Acknowledgments

¹⁰⁴ This work was partially funded by Google Summer of Code 2024 and the Open Source Promotion
¹⁰⁵ Plan (OSPP 2023). We thank Feng Pan for insightful discussions and code contributions on
¹⁰⁶ the slicing technique.

¹⁰⁷ References

- ¹⁰⁸ Bezanson, J., Karpinski, S., Shah, V. B., & Edelman, A. (2012). Julia: A fast dynamic language
¹⁰⁹ for technical computing. *arXiv:1209.5145 [Cs]*. <https://doi.org/10.48550/arXiv.1209.5145>
- ¹¹⁰ Bishop, C. M., & Nasrabadi, N. M. (2006). *Pattern recognition and machine learning* (Vol.
¹¹¹ 4). Springer.
- ¹¹² Bouchitté, V., & Todinca, I. (2001). Treewidth and minimum fill-in: Grouping the minimal
¹¹³ separators. *SIAM Journal on Computing*, 31(1), 212–232.
- ¹¹⁴ Fomin, F. V., & Høie, K. (2006). Pathwidth of cubic graphs and exact algorithms. *Information
¹¹⁵ Processing Letters*, 97(5), 191–196. <https://doi.org/10.1016/j.ipl.2005.10.012>
- ¹¹⁶ Gray, J., & Kourtis, S. (2021). Hyper-optimized tensor network contraction. *Quantum*, 5, 410.
¹¹⁷ <https://doi.org/10.22331/q-2021-03-15-410>
- ¹¹⁸ Haegeman, J., Lubich, C., Oseledets, I., Vandereycken, B., & Verstraete, F. (2016). Unifying
¹¹⁹ time evolution and optimization with matrix product states. *Physical Review B*. <https://doi.org/10.1103/PhysRevB.94.165116>
- ¹²⁰ Harris, C. R., Millman, K. J., van der Walt, S. J., Gommers, R., Virtanen, P., Cournapeau,
¹²¹ D., Wieser, E., Taylor, J., Berg, S., Smith, N. J., Kern, R., Picus, M., Hoyer, S., van
¹²² Kerkwijk, M. H., Brett, M., Haldane, A., del Río, J. F., Wiebe, M., Peterson, P., ...
¹²³ Oliphant, T. E. (2020). Array programming with NumPy. *Nature*, 585(7825), 357–362.
¹²⁴ <https://doi.org/10.1038/s41586-020-2649-2>
- ¹²⁵ Kalachev, G., Panteleev, P., & Yung, M.-H. (2021). *Recursive multi-tensor contraction for
¹²⁶ XEB verification of quantum circuits*. <https://arxiv.org/abs/2108.05665>
- ¹²⁷ Liu, J.-G., Gao, X., Cain, M., Lukin, M. D., & Wang, S.-T. (2023). Computing solution
¹²⁸ space properties of combinatorial optimization problems via generic tensor networks. *SIAM
¹²⁹ Journal on Scientific Computing*, 45(3), A1239–A1270.
- ¹³⁰ Luo, X.-Z., Liu, J.-G., Zhang, P., & Wang, L. (2020). Yao. JI: Extensible, efficient framework
¹³¹ for quantum algorithm design. *Quantum*, 4, 341.
- ¹³² Magron, V., & Wang, J. (2021). TSSOS: A julia library to exploit sparsity for large-scale
¹³³ polynomial optimization. *arXiv:2103.00915*.
- ¹³⁴ Markov, I. L., & Shi, Y. (2008). Simulating Quantum Computation by Contracting Tensor
¹³⁵ Networks. *SIAM Journal on Computing*, 38(3), 963–981. <https://doi.org/10.1137/050644756>
- ¹³⁶ Piveteau, C., Chubb, C. T., & Renes, J. M. (2024). Tensor-Network Decoding Beyond 2D.
¹³⁷ *PRX Quantum*, 5(4), 040303. <https://doi.org/10.1103/PRXQuantum.5.040303>

- ¹⁴⁰ Qing, Y., Li, K., Zhou, P.-F., & Ran, S.-J. (2024). *Compressing neural network by tensor network with exponentially fewer variational parameters* (No. arXiv:2305.06058). arXiv. <https://doi.org/10.48550/arXiv.2305.06058>
- ¹⁴¹ Roa-Villegas, M., Gao, X., Stuijk, S., Corporaal, H., & Liu, J.-G. (2024). Probabilistic Inference in the Era of Tensor Networks and Differential Programming. *Physical Review Research*, 6(3), 033261. <https://doi.org/10.1103/PhysRevResearch.6.033261>
- ¹⁴² Roa-Villegas, M., & Liu, J.-G. (2023). TensorInference: A julia package for tensor-based probabilistic inference. *Journal of Open Source Software*, 8(90), 5700.
- ¹⁴³ Villegas, M. R., Liu, J.-G., Wijnings, P. W. A., Stuijk, S., & Corporaal, H. (2023). Scaling Probabilistic Inference Through Message Contraction Optimization. *2023 Congress in Computer Science, Computer Engineering, & Applied Computing (CSCE)*, 123–130. <https://doi.org/10.1109/CSCE60160.2023.00025>
- ¹⁴⁴
- ¹⁴⁵
- ¹⁴⁶
- ¹⁴⁷
- ¹⁴⁸
- ¹⁴⁹
- ¹⁵⁰
- ¹⁵¹