

¹ OMEinsumContractionOrders: A Julia package for tensor network contraction order optimization

³ Jin-Guo Liu  ^{1*}, Xuanzhao Gao^{2*}, and Richard Samuelson^{3*}

⁴ 1 Hong Kong University of Science and Technology (Guangzhou) 2 Center of Computational
⁵ Mathematics, Flatiron Institute 3 * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))

⁶ Summary

⁷ OMEinsumContractionOrders (One More Einsum Contraction Orders, or OMECO) is a Julia package ([Bezanson et al., 2012](#)) that implements state-of-the-art algorithms for optimizing ⁸ tensor network contraction orders. OMECO is designed to search for near-optimal contraction ⁹ orders for exact tensor network contraction, and provides a comprehensive suite of optimization ¹⁰ algorithms for tensor network contraction orders, including greedy heuristics, simulated ¹¹ annealing, and tree width solvers. In this paper, we present the key features of OMECO, its ¹² integration with the Julia ecosystem, and performance benchmarks.

¹⁴ Statement of need

A *tensor network* is a mathematical framework that represents multilinear algebra operations as graphical structures, where tensors are nodes and shared indices are edges. This diagrammatic approach transforms complex high-dimensional contractions into visual networks that expose underlying computational structure.

¹⁹ The framework has remarkable universality across diverse domains: *einsum* notation ([Harris et al., 2020](#)) in numerical computing, *factor graphs* ([Bishop & Nasrabadi, 2006](#)) in probabilistic inference, *sum-product networks* in machine learning, and *junction trees* ([Villegas et al., 2023](#)) in graphical models. Tensor networks have enabled breakthroughs in quantum circuit simulation ([Markov & Shi, 2008](#)), quantum error correction ([Piveteau et al., 2024](#)), neural network compression ([Qing et al., 2024](#)), strongly correlated quantum materials ([Haegeman et al., 2016](#)), and combinatorial optimization problems ([J.-G. Liu et al., 2023](#)).

²⁶ The computational cost of tensor network contraction depends critically on the *contraction order*—the sequence in which pairwise tensor multiplications are performed. This order ²⁷ can be represented as a binary tree where leaves correspond to input tensors and internal ²⁸ nodes represent intermediate results. The optimization objective balances multiple complexity ²⁹ measures through the cost function:

$$\mathcal{L} = w_t \cdot tc + w_s \cdot \max(0, sc - sc_{target}) + w_{rw} \cdot rwc,$$

³¹ where w_t , w_s , and w_{rw} represent weights for time complexity (tc), space complexity (sc), and ³² read-write complexity (rwc), respectively. In practice, memory access costs typically dominate ³³ computational costs, motivating $w_{rw} > w_t$. The space complexity penalty activates only when ³⁴ $sc > sc_{target}$, allowing unconstrained optimization when memory fits within available device ³⁵ capacity.

³⁶ Finding the optimal contraction order—even when minimizing only time complexity—is NP-³⁷ complete ([Markov & Shi, 2008](#)). Algorithms for finding near-optimal contraction orders ³⁸ have been developed and achieve impressive scalability ([Gray & Kourtis, 2021](#); [Roa-Villegas](#)

³⁹ et al., 2024), handling tensor networks with over 10^3 tensors. While the Python package ⁴⁰ cotengra (Gray & Kourtis, 2021) has been widely adopted in the community, achieving optimal ⁴¹ performance across diverse problem instances—particularly when balancing solution quality ⁴² against optimization time constraints—remains an open challenge.

⁴³ OMECO addresses this challenge through a unified and extensible framework that integrates ⁴⁴ multiple complementary optimization strategies, including greedy heuristics, simulated ⁴⁵ annealing, and tree-width-based solvers. This comprehensive approach enables more systematic ⁴⁶ exploration of the optimization time-solution quality trade-off space. OMECO has been in- ⁴⁷ tegrated into the OMEinsum package and powers several downstream applications: Yao (Luo ⁴⁸ et al., 2020) for quantum circuit simulation, GenericTensorNetworks (J.-G. Liu et al., 2023) ⁴⁹ and TensorBranching for combinatorial optimization, TensorInference (Roa-Villegas & Liu, ⁵⁰ 2023) for probabilistic inference, and TensorQEC for quantum error correction. This infrastruc- ⁵¹ ture is expected to benefit other applications requiring tree or path decomposition, such as ⁵² polynomial optimization (Magron & Wang, 2021). These applications are reflected in the ⁵³ ecosystem built around OMECO, as illustrated in Figure 1.

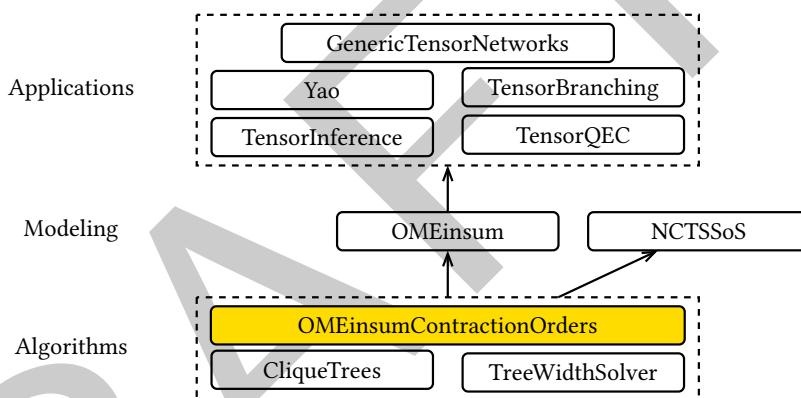


Figure 1: The ecosystem built around OMEinsumContractionOrders and its dependencies. OMECO serves as a core component of the tensor network contractor OMEinsum, which powers applications including Yao (quantum simulation), TensorQEC (quantum error correction), TensorInference (probabilistic inference), GenericTensorNetworks and TensorBranching (combinatorial optimization).

54 Features and benchmarks

⁵⁵ The major feature of OMECO is contraction order optimization. OMECO provides several ⁵⁶ algorithms with complementary performance characteristics that can be simply called by the ⁵⁷ optimize_code function:

Optimizer	Description
GreedyMethod	Fast greedy heuristic with modest solution quality
TreeSA	Reliable simulated annealing optimizer (Kalachev et al., 2021) with high-quality solutions
PathSA	Simulated annealing optimizer for path decomposition
HyperND	Nested dissection algorithm for hypergraphs, requires KaHyPar or Metis
KaHyParBipartite	Graph bipartition method for large tensor networks (Gray & Kourtis, 2021), requires KaHyPar

Optimizer	Description
SABipartite	Simulated annealing bipartition method, pure Julia implementation
ExactTreewidth	Exact algorithm with exponential runtime (Bouchitté & Todinca, 2001), based on TreeWidthSolver
Treewidth	Clique tree elimination methods from CliqueTrees package (Samuelson & Fairbanks, 2025)

58 The algorithms HyperND, Treewidth, and ExactTreewidth are tree-width based solvers that
59 operate on graphs. They first convert tensor networks to their line graph representation ([Markov](#)
60 & [Shi, 2008](#)) and then find an optimized tree decomposition of the line graph using the
61 CliqueTrees and TreeWidthSolver packages, as illustrated in [Figure 1](#). Additionally, the
62 PathSA optimizer optimizes path decomposition instead of tree decomposition. It is a variant
63 of TreeSA by constraining contraction orders to path graphs, which is useful for applications
64 requiring a linear contraction order.

65 These methods balance optimization time against solution quality. [Figure 2](#) displays benchmark
66 results for the tensor network of the Sycamore quantum circuit ([Arute et al., 2019; Pan &](#)
67 [Zhang, 2021](#)) that widely used as a benchmark for quantum supremacy, which is believed to
68 have an optimal space complexity of 52. The Pareto front highlights the optimal trade-off
69 between optimization time and solution quality.

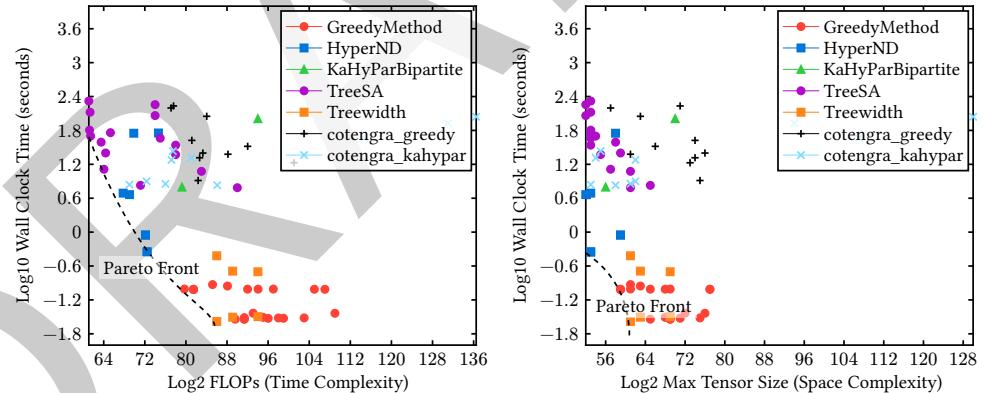


Figure 2: Time complexity (left) and space complexity (right) benchmark results for contraction order optimization on the Sycamore quantum circuit tensor network (Intel Xeon Gold 6226R CPU @ 2.90GHz, single-threaded). The x -axis shows contraction cost, y -axis shows optimization time. Each point represents a different optimizer configuration tested with varying parameters. TreeSA and HyperND achieve the lowest contraction costs, while GreedyMethod offers the fastest optimization time. The parameter setup for each optimizer is detailed in our benchmark repository [OMEInsumContractionOrdersBenchmark](#).

70 Optimizers prefixed with `cotengra_` are from the Python package `cotengra` ([Gray & Kourtis,](#)
71 [2021](#)); all others are OMECO implementations. For both optimization objectives (minimizing
72 time and space complexity), OMECO optimizers dominate the Pareto front. Given suffi-
73 cient optimization time, TreeSA consistently achieves the lowest time and space complexity.
74 GreedyMethod and Treewidth (backed by minimum fill (MF) ([Ng & Peyton, 2014](#)), multiple
75 minimum degree (MMD) ([J. W. Liu, 1985](#)), and approximate minimum fill (AMF) ([Rothberg](#)
76 & [Eisenstat, 1998](#))) provides the fastest optimization but yields suboptimal contraction orders,
77 while HyperND offers a favorable balance between optimization time and solution quality.

78 More real-world examples demonstrating applications to quantum circuit simulation, combinatorial optimization, and probabilistic inference are available in the [OMEinsumContractionOrder-
sBenchmark](#) repository. We find that optimizer performance is highly problem-dependent, with
81 no single algorithm dominating across all metrics and graph topologies.

82 Another key feature of OMECO is index slicing, a technique that trades time complexity for
83 reduced space complexity by explicitly looping over a subset of tensor indices. OMECO provides
84 the `slice_code` interface for this purpose, currently supporting the TreeSASlicer algorithm,
85 which implements dynamic slicing based on the TreeSA optimizer. [Figure 3](#) demonstrates this
86 capability using the Sycamore quantum circuit, where slicing reduces the space complexity
87 from 2^{52} to 2^{31} .

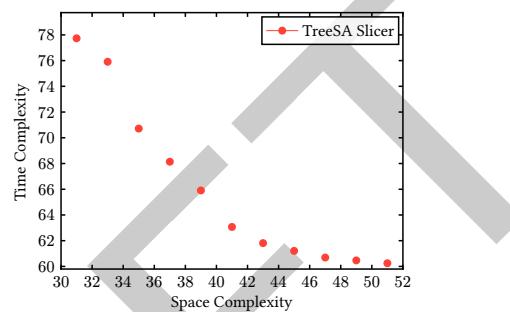


Figure 3: Trade-off between time complexity and target space complexity using TreeSASlicer on the Sycamore quantum circuit. The original network has a space complexity of 2^{52} .

88 The numerical experiments show that moderate slicing increases time complexity only slightly,
89 while aggressive slicing can induce significant overhead. There is a critical transition point
90 around 42 where the time complexity begins to increase significantly.
91 We thank the Julia community and all contributors to the `OMEinsum` and `OMEinsumContractionOrders`
92 packages. We are grateful to Xiwei Pan for valuable writing suggestions that improved this
93 manuscript.

References

- 95 Arute, F., Arya, K., Babbush, R., Bacon, D., Bardin, J. C., Barends, R., Biswas, R., Boixo, S.,
96 Branda, F. G., Buell, D. A., & others. (2019). Quantum supremacy using a programmable
97 superconducting processor. *Nature*, 574(7779), 505–510. <https://www.nature.com/articles/s41586-019-1666-5>
- 99 Bezanson, J., Karpinski, S., Shah, V. B., & Edelman, A. (2012). Julia: A fast dynamic language
100 for technical computing. *arXiv:1209.5145 [Cs]*. <https://doi.org/10.48550/arXiv.1209.5145>
- 101 Bishop, C. M., & Nasrabadi, N. M. (2006). *Pattern recognition and machine learning* (Vol.
102 4). Springer.
- 103 Bouchitté, V., & Todinca, I. (2001). Treewidth and minimum fill-in: Grouping the minimal
104 separators. *SIAM Journal on Computing*, 31(1), 212–232. <https://doi.org/10.1137/S0097539799359683>
- 106 Gray, J., & Kourtis, S. (2021). Hyper-optimized tensor network contraction. *Quantum*, 5, 410.
107 <https://doi.org/10.22331/q-2021-03-15-410>
- 108 Haegeman, J., Lubich, C., Oseledets, I., Vandereycken, B., & Verstraete, F. (2016). Unifying
109 time evolution and optimization with matrix product states. *Physical Review B*. <https://doi.org/10.1103/PhysRevB.94.165116>

- 111 Harris, C. R., Millman, K. J., van der Walt, S. J., Gommers, R., Virtanen, P., Cournapeau,
 112 D., Wieser, E., Taylor, J., Berg, S., Smith, N. J., Kern, R., Picus, M., Hoyer, S., van
 113 Kerkwijk, M. H., Brett, M., Haldane, A., del Río, J. F., Wiebe, M., Peterson, P., ...
 114 Oliphant, T. E. (2020). Array programming with NumPy. *Nature*, 585(7825), 357–362.
 115 <https://doi.org/10.1038/s41586-020-2649-2>
- 116 Kalachev, G., Panteleev, P., & Yung, M.-H. (2021). *Recursive multi-tensor contraction for*
 117 *XEB verification of quantum circuits*. <https://arxiv.org/abs/2108.05665>
- 118 Liu, J. W. (1985). Modification of the minimum-degree algorithm by multiple elimination.
 119 *ACM Transactions on Mathematical Software (TOMS)*, 11(2), 141–153. <https://doi.org/10.1145/214392.214398>
- 120 Liu, J.-G., Gao, X., Cain, M., Lukin, M. D., & Wang, S.-T. (2023). Computing solution space
 121 properties of combinatorial optimization problems via generic tensor networks. *SIAM Journal*
 122 *on Scientific Computing*, 45(3), A1239–A1270. <https://doi.org/10.1137/22M1501787>
- 123 Luo, X.-Z., Liu, J.-G., Zhang, P., & Wang, L. (2020). Yao. Jl: Extensible, efficient
 124 framework for quantum algorithm design. *Quantum*, 4, 341. <https://doi.org/10.22331/q-2020-10-11-341>
- 125 Magron, V., & Wang, J. (2021). TSSOS: A julia library to exploit sparsity for large-scale
 126 polynomial optimization. *arXiv:2103.00915*. <https://arxiv.org/abs/2103.00915>
- 127 Markov, I. L., & Shi, Y. (2008). Simulating Quantum Computation by Contracting Tensor
 128 Networks. *SIAM Journal on Computing*, 38(3), 963–981. <https://doi.org/10.1137/050644756>
- 129 Ng, E. G., & Peyton, B. W. (2014). Fast implementation of the minimum local fill ordering
 130 heuristic. *CSC14: The Sixth SIAM Workshop on Combinatorial Scientific Computing*, 4.
 131 <https://doi.org/10.1137/070680680>
- 132 Pan, F., & Zhang, P. (2021). *Simulating the sycamore quantum supremacy circuits*. <https://arxiv.org/abs/2103.03074>
- 133 Piveteau, C., Chubb, C. T., & Renes, J. M. (2024). Tensor-Network Decoding Beyond 2D.
 134 *PRX Quantum*, 5(4), 040303. <https://doi.org/10.1103/PRXQuantum.5.040303>
- 135 Qing, Y., Li, K., Zhou, P.-F., & Ran, S.-J. (2024). *Compressing neural network by tensor*
 136 *network with exponentially fewer variational parameters* (No. arXiv:2305.06058). arXiv.
 137 <https://doi.org/10.48550/arXiv.2305.06058>
- 138 Roa-Villegas, M., Gao, X., Stuijk, S., Corporaal, H., & Liu, J.-G. (2024). Probabilistic
 139 Inference in the Era of Tensor Networks and Differential Programming. *Physical Review*
 140 *Research*, 6(3), 033261. <https://doi.org/10.1103/PhysRevResearch.6.033261>
- 141 Roa-Villegas, M., & Liu, J.-G. (2023). TensorInference: A julia package for tensor-based
 142 probabilistic inference. *Journal of Open Source Software*, 8(90), 5700. <https://joss.theoj.org/papers/10.21105/joss.05700>
- 143 Rothberg, E., & Eisenstat, S. C. (1998). Node selection strategies for bottom-up sparse
 144 matrix ordering. *SIAM Journal on Matrix Analysis and Applications*, 19(3), 682–695.
 145 <https://doi.org/10.1137/S0895479896302692>
- 146 Samuelson, R., & Fairbanks, J. (2025). *CliqueTrees.jl: A julia library for computing tree*
 147 *decompositions and chordal completions of graphs*. <https://github.com/AlgebraicJulia/CliqueTrees.jl>
- 148 Villegas, M. R., Liu, J.-G., Wijnings, P. W. A., Stuijk, S., & Corporaal, H. (2023). Scaling
 149 Probabilistic Inference Through Message Contraction Optimization. *2023 Congress in*
 150 *Computer Science, Computer Engineering, & Applied Computing (CSCE)*, 123–130. <https://doi.org/10.1109/CSCE60160.2023.00025>